

Diabetes prediction using different classification models, calculating accuracy of each model and finding the model which gives high accuracy

Diabetes is a type of chronic disease which is more common among the people of all age groups. Predicting this disease at an early stage can help a person to take the necessary precautions and change his/her lifestyle accordingly to either prevent the occurrence of this disease or control the disease(For people who already have the disease).

Steps:

- > Loading Dataset
- > Visualizing and cleaning Dataset
- > Splitting the dataset into training set and testing set
- > Building different classification models, predicting values and accuracy
- > Finding the highest accuracy

LOADING DATASET

Importing libraries pandas, numpy and matplotlib

```
In [14]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

Reading data set , these dataset is regarding the features of different people and outcome in the form of binary number 1 or 0. 1 - indicates diabetes is present, 0- indicates diabetes is absent. This dataset is downloaded from UCI machine learning repository

```
In [15]: df = pd.read_csv(r'Documents\Machine-Learning-with-Python-master\diabetes.csv')
```

Analyzing and Visualizing Dataset

```
In [4]: df.head()
```

```
Out[4]:
```

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age |
|---|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | 0.627 | 51 |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | 0.351 | 33 |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | 0.672 | 33 |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | 0.167 | 40 |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | 2.288 | 33 |

We can observe from the above data set it has different features like Pregnancies, Glucose, BloodPressure, Skin Thickness, Insulin , BMI , DiabetesPedigreeFunction, Age. Here label is Outcome

```
In [5]: print("shape of data set is {}".format(df.shape))
```

```
shape of data set is (768, 9)
```

The above data set has total 768 rows and 9 columns


```
In [6]: df.groupby('Outcome').size()
```

```
Out[6]: Outcome
0      500
1      268
dtype: int64
```

It has total 268 Diabetes positive cases and 500 diabetes negative cases

```
In [26]: import seaborn as sns
sns.countplot(df['Outcome'], label='count')
```

```
Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x1f665eaf748>
```



STATISTICAL MEASUREMENTS

we can also calculate different statistical measurements like mean, standard deviation of each feature of a data set

In [5]: `df.describe()`

Out[5]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesP |
|--------------|-------------|------------|---------------|---------------|------------|------------|-----------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | |
| mean | 3.845052 | 120.894531 | 69.105469 | 20.536458 | 79.799479 | 31.992578 | |
| std | 3.369578 | 31.972618 | 19.355807 | 15.952218 | 115.244002 | 7.884160 | |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 25% | 1.000000 | 99.000000 | 62.000000 | 0.000000 | 0.000000 | 27.300000 | |
| 50% | 3.000000 | 117.000000 | 72.000000 | 23.000000 | 30.500000 | 32.000000 | |
| 75% | 6.000000 | 140.250000 | 80.000000 | 32.000000 | 127.250000 | 36.600000 | |
| max | 17.000000 | 199.000000 | 122.000000 | 99.000000 | 846.000000 | 67.100000 | |

We can get data types of different features by using the below Data code

In [21]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
Pregnancies      768 non-null int64
Glucose          768 non-null int64
BloodPressure    768 non-null int64
SkinThickness    768 non-null int64
Insulin          768 non-null int64
BMI              768 non-null float64
DiabetesPedigreeFunction  768 non-null float64
Age              768 non-null int64
Outcome          768 non-null int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

SPLITTING THE DATASET INTO TRAINING AND TESTING SETS

TRAINING SET : The part of the dataset which we are going to train to the model

TESTING SET : The part of the dataset on which we are going to test out model

In order to split the data into training and testing data set we should import the `train_test_split` from the `sklearn` module

In [16]: `from sklearn.model_selection import train_test_split`
`x_train,x_test,y_train,y_test = train_test_split(df.loc[:,df.columns!='Outcome'],df['Outcome'],stratify=df['Outcome'],random_state=66)`

Features which we are going to train

In [9]: `x_train.head()`

Out[9]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|
| 612 | 7 | 168 | 88 | 42 | 321 | 38.2 | 0.787 |
| 557 | 8 | 110 | 76 | 0 | 0 | 27.8 | 0.237 |
| 26 | 7 | 147 | 76 | 0 | 0 | 39.4 | 0.257 |
| 70 | 2 | 100 | 66 | 20 | 90 | 32.9 | 0.867 |
| 73 | 4 | 129 | 86 | 20 | 270 | 35.1 | 0.231 |

Features set of the testing data

```
In [8]: x_test.head()
```

Out[8]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|
| 506 | 0 | 180 | 90 | 26 | 90 | 36.5 | 0.314 |
| 709 | 2 | 93 | 64 | 32 | 160 | 38.0 | 0.674 |
| 257 | 2 | 114 | 68 | 22 | 0 | 28.7 | 0.092 |
| 518 | 13 | 76 | 60 | 0 | 0 | 32.8 | 0.180 |
| 432 | 1 | 80 | 74 | 11 | 60 | 30.0 | 0.527 |

Label set of the Training part of the dataset

```
In [9]: y_train.head()
```

Out[9]:

| | |
|-----|---|
| 612 | 1 |
| 557 | 0 |
| 26 | 1 |
| 70 | 1 |
| 73 | 0 |

Name: Outcome, dtype: int64

Label set of the Testing part of the dataset

```
In [10]: y_test.head()
```

Out[10]:

| | |
|-----|---|
| 506 | 1 |
| 709 | 1 |
| 257 | 0 |
| 518 | 0 |
| 432 | 0 |

Name: Outcome, dtype: int64

In this project the classification models used are:

- 1) k Nearest neighbors model
- 2) Logistic regression model

- 3) Decision Trees
- 4) Support Vector Machine

Building model using k-nearest neighbors classification

The k-NN algorithm is the simplest machine learning algorithm. Building the model consists only of storing the training data set. To make a prediction for a new data point, the algorithm finds the closest data points in the training data set — its “nearest neighbors.”

Let us predict the outcome with different values of k (no_of_neighbors)

```
In [17]: # K-nearest-neighbors classifier
list_of_training_accuracy = []
list_of_testing_accuracy = []
from sklearn.neighbors import KNeighborsClassifier
training_accuracy = []
testing_accuracy = []
neighbors = list(range(1,10))
for no_of_neighbors in neighbors:
    kn = KNeighborsClassifier(n_neighbors=no_of_neighbors).fit(x_train,y_train)
    y_pre = kn.predict(x_test)
    print("predicted values of k = {} is ".format(no_of_neighbors))
    print(y_pre)
```

```

predicted values of k = 1 is
[1 0 0 0 0 0 0 1 0 0 0 1 0 1 0 1 0 0 0 1 1 0 0 1 0 0 0 0 1 1 1 1 1 0 0 0 0
0
0 0 1 1 0 1 0 1 0 0 1 0 1 0 1 0 0 1 1 1 1 0 1 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0
0
0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 1 0 0 1 0 0 0 0 0 1 1 1 0 1 1 1 0 0 1
1
0 1 0 0 0 0 1 0 1 1 0 0 1 1 0 1 0 1 0 0 0 1 0 0 0 1 1 0 0 1 1 1 1 1 0 0
0
1 1 1 0 1 0 0 0 1 0 0 1 0 0 1 1 0 0 0 0 1 1 0 1 1 0 0 1 0 0 1 1 0 0 1 1
0
0 0 1 0 0 0 0]
predicted values of k = 2 is
[1 0 0 0 0 0 0 1 0 0 0 1 0 1 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0
0 0 0 1 0 0 0 1 0 0 1 0 1 0 1 0 0 1 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0

```

```

0
0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0
1
0 1 0 0 0 0 1 0 1 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 1 1 0 0
0
0 1 0 0 1 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 1
0
0 0 1 0 0 0 0]
predicted values of k = 3 is
[1 0 0 0 0 1 1 0 0 1 1 0 1 0 0 0 0 1 1 1 1 0 0 0 0 0 1 1 0 0 1 1 0 0 0 1
0
0 1 0 1 0 0 0 1 0 0 1 0 1 0 1 0 0 1 1 1 1 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0
0
0 0 0 1 0 0 0 1 1 0 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 0 1 1 1 0 0 1 1 1 0 0
1
0 1 0 0 1 0 1 0 1 0 0 0 0 1 1 1 0 1 0 0 0 0 1 0 0 0 1 1 0 0 1 1 1 1 0 0
0
0 1 1 0 1 1 0 1 0 1 0 1 0 0 0 1 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 1 1 0 1 0 1
0
1 0 1 0 1 0 0]
predicted values of k = 4 is
[1 0 0 0 0 1 1 0 0 1 1 0 1 0 0 0 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0
0 1 0 1 0 0 0 1 0 0 1 0 1 0 1 0 0 1 1 0 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0
0
0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 1 1 0 0 0 0 1 1 0 0 0 0 1 0 0 0
0
0 1 0 0 0 0 1 0 0 0 0 0 0 1 1 1 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 1 1 0 0
0
0 1 0 0 1 0 0 1 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 1
0
1 0 1 0 1 0 0]
predicted values of k = 5 is
[1 0 0 0 0 1 1 1 0 1 1 0 1 0 0 0 0 1 1 1 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0
1
1 1 0 1 0 0 0 1 0 0 1 0 1 0 1 0 0 1 1 1 0 0 1 0 1 0 0 1 0 0 0 0 0 0 0 0
0
0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 1 1 0 0 1 1 0 0 0 0 1 1 0 0 0 0 1 0 0 0
1
0 1 0 0 0 0 1 0 0 0 0 0 0 1 1 1 0 1 0 0 0 0 1 0 0 0 1 1 0 0 1 1 1 1 0 0
0
1 1 0 0 1 0 0 1 0 0 0 1 1 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 1
0
1 0 1 0 1 1 0]
predicted values of k = 6 is
[0 0 0 0 0 1 1 0 0 0 1 0 1 0 0 0 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0

```


Let us calculate the accuracy of each model for each value of K which are used above

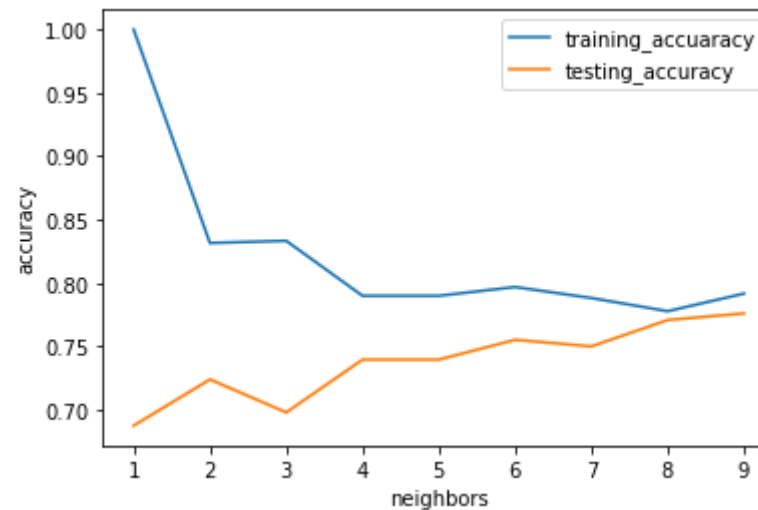
```
In [18]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
for no_of_neighbors in neighbors:
    kn = KNeighborsClassifier(n_neighbors=no_of_neighbors).fit(x_train,y_train)
    training_accuracy.append(kn.score(x_train,y_train))
    testing_accuracy.append(kn.score(x_test,y_test))
print("list of training accuracy of differnt k values models")
print(training_accuracy)
print('\n')
print('list of testing accuracy of differnt k values models')
print(testing_accuracy)
```

```
list of training accuracy of differnt k values models
[1.0, 0.8315972222222222, 0.8333333333333334, 0.7899305555555556, 0.7899305555555556, 0.796875, 0.7881944444444444, 0.7777777777777778, 0.7916666666666666]
```

```
list of testing accuracy of differnt k values models
[0.6875, 0.7239583333333334, 0.6979166666666666, 0.7395833333333334, 0.7395833333333334, 0.7552083333333334, 0.75, 0.7708333333333334, 0.7760416666666666]
```

Plotting the graph of between accuracies of training and testing datasets for differnt values of k

```
In [19]: plt.plot(neighbors,training_accuracy,label="training_accuaracy")
plt.plot(neighbors,testing_accuracy,label="testing_accuracy")
plt.ylabel('accuracy')
plt.xlabel('neighbors')
plt.legend()
plt.savefig('knn_accuracy_comapare_model')
```



from the above graph we can observe that we should we are accuracy of training and testing dataset at equal level at around 8-9, hence we can consider that k value around 9

Hence for k equal to around 8-9 we are getting the matching accuracy for both training and testing sets. The model is built with the value of $k = 9$

Confusion Matrix: A confusion matrix is a summary of prediction results on a classification problem. The number of correct and incorrect predictions are summarized with count values and broken down by each class. This is the key to the confusion matrix. The confusion matrix shows the ways in which your classification model is confused when it makes predictions. It gives us insight not only into the errors being made by a classifier but more importantly the types of errors that are being made.

| | <i>Class 1 Predicted</i> | <i>Class 2 Predicted</i> |
|---------------------------|------------------------------|------------------------------|
| Class 1 Actual | TP | FN |
| Class 2 Actual | FP | TN |

Here, Class 1 : Positive Class 2 : Negative

Definition of the Terms:

Positive (P) : Observation is positive (for example: is an apple). Negative (N) : Observation is not positive (for example: is not an apple). True Positive (TP) : Observation is positive, and is predicted to be positive. False Negative (FN) : Observation is positive, but is predicted negative. True Negative (TN) : Observation is negative, and is predicted to be negative. False Positive (FP) : Observation is negative, but is predicted positive

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$

High recall, low precision: This means that most of the positive examples are correctly recognized (low FN) but there are a lot of false positives.

Low recall, high precision: This shows that we miss a lot of positive examples (high FN) but those we predict as positive are indeed positive (low FP)

F-measure: Since we have two measures (Precision and Recall) it helps to have a measurement that represents both of them. We calculate an F-measure which uses Harmonic Mean in place of Arithmetic Mean as it punishes the extreme values more. The F-Measure will always be nearer to the smaller value of Precision or Recall.

$$F - measure = \frac{2 * Recall * Precision}{Recall + Precision}$$

```
In [20]: # from the above figure we can say that we have to neighbor some where a
round 9
kn = KNeighborsClassifier(n_neighbors = 9).fit(x_train,y_train)
list_of_training_accuracy.append(kn.score(x_train,y_train))
list_of_testing_accuracy.append(kn.score(x_test,y_test))
print("Accuracy of the K-Neares_neighbors_classifer on training set is "
, kn.score(x_train,y_train))
print("Accuracy of the K-Neares_neighbors_classifer on testing set is ",
kn.score(x_test,y_test))
print()
print("Confusion Matrix ")
print(confusion_matrix(y_test, kn.predict(x_test)))
print()
print("Report")
print(classification_report(y_test, kn.predict(x_test)))
```

```
Accuracy of the K-Neares_neighbors_classifer on training set is  0.79166
666666666666
```

```
Accuracy of the K-Neares_neighbors_classifer on testing set is  0.776041
666666666666
```

```
Confusion Matrix
```

```
[[105  20]
 [ 23  44]]
```

| Report | | precision | recall | f1-score | support |
|--------|--------------|-----------|--------|----------|---------|
| | 0 | 0.82 | 0.84 | 0.83 | 125 |
| | 1 | 0.69 | 0.66 | 0.67 | 67 |
| | accuracy | | | 0.78 | 192 |
| | macro avg | 0.75 | 0.75 | 0.75 | 192 |
| | weighted avg | 0.77 | 0.78 | 0.77 | 192 |

Building model using logistic regression

This type of statistical analysis (also known as logit model) is often used for predictive analytics and modeling, and extends to applications in machine learning. In this analytics approach, the dependent variable is finite or categorical: either A or B (binary regression) or a range of finite options A, B, C or D (multinomial regression). It is used in statistical software to understand the relationship between the dependent variable and one or more independent variables by estimating probabilities using a logistic regression equation.

This type of analysis can help you predict the likelihood of an event happening or a choice being made. For example, you may want to know the likelihood of a visitor choosing an offer made on your website — or not (dependent variable). Your analysis can look at known characteristics of visitors, such as sites they came from, repeat visits to your site, behavior on your site (independent variables). Logistic regression models help you determine a probability of what type of visitors are likely to accept the offer — or not. As a result, you can make better decisions about promoting your offer or make decisions about the offer itself.

```
In [21]: # logistic Regression
from sklearn.linear_model import LogisticRegression
lgre = LogisticRegression().fit(x_train,y_train)
ypre = lgre.predict(x_test)
print("predicted values",ypre)
print("test set values",list(y_test))
list_of_training_accuracy.append(lgre.score(x_train,y_train))
list_of_testing_accuracy.append(lgre.score(x_test,y_test))
print()
print("Accuracy of Logistic_Regression on training set is ",lgre.score(x_train,y_train))
print("Accuracy of Logistic_Regression on testing set is ",lgre.score(x_test,y_test))
```

```

print()
print("Confusion Matrix ")
print(confusion_matrix(y_test,lgre.predict(x_test)))
print()
print("Report")
print(classification_report(y_test,lgre.predict(x_test)))

```

C:\anaconda\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
FutureWarning)

```

predicted values [1 0 0 0 0 1 1 0 0 0 1 0 1 0 1 0 0 1 1 1 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0
1 1 0 1 0 0 0 1 0 0 1 0 1 1 1 0 0 1 0 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0
0
0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 1 1 1 0 1 0 0 0 0 0 1 0 0 0 0 0 1 0 0
1
0 1 0 0 1 0 1 0 0 0 0 0 1 0 1 1 0 1 0 0 1 0 0 0 0 0 0 1 0 0 1 1 1 1 0 0
1
0 0 0 0 1 0 0 1 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 1 1
1
0 0 0 0 1 0 0]
test set values [1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1,
1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0,
0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0,
0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0,
0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,
0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1,
1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0,
0, 1, 0, 1, 1, 0]

```

Accuracy of Logistic_Regression on training set is 0.78125
Accuracy of Logistic_Regression on testing set is 0.7708333333333334

Confusion Matrix
[[110 15]
[29 38]]

Report

| | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.79 | 0.88 | 0.83 | 125 |
| 1 | 0.72 | 0.57 | 0.63 | 67 |

| | | | | |
|--------------|------|------|------|-----|
| accuracy | | | 0.77 | 192 |
| macro avg | 0.75 | 0.72 | 0.73 | 192 |
| weighted avg | 0.77 | 0.77 | 0.76 | 192 |

Building a model using Decision tree

Decision tree algorithm falls under the category of supervised learning. They can be used to solve both regression and classification problems. Decision tree uses the tree representation to solve the problem in which each leaf node corresponds to a class label and attributes are represented on the internal node of the tree. We can represent any boolean function on discrete attributes using the decision tree.

```
In [22]: #decision trees

from sklearn.tree import DecisionTreeClassifier
dtree = DecisionTreeClassifier(max_depth = 3).fit(x_train,y_train)
ypre=dtree.predict(x_test)
print("Predicted values of Outcome using Decison Trees classfier")
print(ypre)
print()
print("actual values of Outcome of the order set")
print(list(y_test))
list_of_training_accuracy.append(dtree.score(x_train,y_train))
list_of_testing_accuracy.append(dtree.score(x_test,y_test))
print()
print("Accuracy of Decision Trees on training set is ",dtree.score(x_train,y_train))
print("Accuracy of Decision Trees on testing set is ",dtree.score(x_test,y_test))
print()
print("Confusion Matrix ")
print(confusion_matrix(y_test,kn.predict(x_test)))
print()
print("Report")
print(classification_report(y_test,kn.predict(x_test)))
```

```
Predicted values of Outcome using Decison Trees classfier
[1 0 0 0 0 1 1 0 0 0 1 0 1 0 1 0 0 1 1 1 0 0 0 0 1 0 0 0 0 0 0 1 0 1 0 1 1
1
1 0 0 1 0 0 0 1 0 0 1 0 1 1 1 0 0 0 0 1 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0
0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 1 0 0 1 0 0 0 0 0 1 0 1 0 1 0 0 1 0 0
1
```

```

1
0 1 0 0 1 0 1 0 0 0 1 0 1 0 1 1 0 1 0 0 1 0 0 0 0 1 1 1 1 1 1 1 1 0 0
1
0 0 0 0 1 1 0 1 0 1 0 1 1 1 0 0 0 0 0 0 0 0 0 0 1 1 0 1 0 0 0 0 1 0 0 1 1
0
0 0 0 0 1 0 0]

```

actual values of Outcome of the order set

```

[1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1,
0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1,
0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1,
1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0,
0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1,
0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0]

```

Accuracy of Decision Trees on training set is 0.7725694444444444

Accuracy of Decision Trees on testing set is 0.7395833333333334

Confusion Matrix

```

[[105  20]
 [ 23  44]]

```

Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.82 | 0.84 | 0.83 | 125 |
| 1 | 0.69 | 0.66 | 0.67 | 67 |
| accuracy | | | 0.78 | 192 |
| macro avg | 0.75 | 0.75 | 0.75 | 192 |
| weighted avg | 0.77 | 0.78 | 0.77 | 192 |

Building a model using Support Vector Machine

SVM is a supervised machine learning algorithm which can be used for classification or regression problems. It uses a technique called the kernel trick to transform your data and then based on these transformations it finds an optimal boundary between the possible outputs. Simply put, it does some extremely complex data transformations, then figures out how to separate your data based on the labels or outputs you've defined.


```
In [23]: #support vector Machine
from sklearn.svm import SVC
svc = SVC()
svc.fit(x_train, y_train)
print("Accuracy on training set: {:.2f}".format(svc.score(x_train, y_train)))
print("Accuracy on test set: {:.2f}".format(svc.score(x_test, y_test)))
print()
print("Confusion Matrix ")
print(confusion_matrix(y_test, kn.predict(x_test)))
print()
print("Report")
print(classification_report(y_test, kn.predict(x_test)))
```

Accuracy on training set: 1.00

Accuracy on test set: 0.65

Confusion Matrix

```
[[105  20]
 [ 23  44]]
```

Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.82 | 0.84 | 0.83 | 125 |
| 1 | 0.69 | 0.66 | 0.67 | 67 |
| accuracy | | | 0.78 | 192 |
| macro avg | 0.75 | 0.75 | 0.75 | 192 |
| weighted avg | 0.77 | 0.78 | 0.77 | 192 |

C:\anaconda\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
 "avoid this warning.", FutureWarning)

The accuracy on the training set is 100%, while the test set accuracy is much worse. This is an indicative that the tree is overfitting and not generalizing well to new data. Therefore, we need to apply pre-pruning to the tree. We set max_depth=3, limiting the depth of the tree decreases overfitting. This leads to a lower accuracy on the training set, but an improvement on the test set.

```
In [24]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.fit_transform(x_test)
svc = SVC()
svc.fit(x_train_scaled, y_train)
print("Accuracy on training set: {:.2f}".format(svc.score(x_train_scaled
, y_train)))
print("Accuracy on test set: {:.2f}".format(svc.score(x_test_scaled, y_t
est)))
print()
print("Confusion Matrix ")
print(confusion_matrix(y_test, kn.predict(x_test)))
print()
print("Report")
print(classification_report(y_test, kn.predict(x_test)))
```

Accuracy on training set: 0.77

Accuracy on test set: 0.77

Confusion Matrix

```
[[105  20]
 [ 23  44]]
```

Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.82 | 0.84 | 0.83 | 125 |
| 1 | 0.69 | 0.66 | 0.67 | 67 |
| accuracy | | | 0.78 | 192 |
| macro avg | 0.75 | 0.75 | 0.75 | 192 |
| weighted avg | 0.77 | 0.78 | 0.77 | 192 |

C:\anaconda\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
 "avoid this warning.", FutureWarning)

```
In [25]: svc = SVC(C=1000)
svc.fit(x_train_scaled, y_train)
list_of_training_accuracy.append(svc.score(x_train_scaled,y_train))
```

```

list_of_testing_accuracy.append(svc.score(x_test_scaled,y_test))
print("Accuracy on training set: {:.3f}".format(
    svc.score(x_train_scaled, y_train)))
print("Accuracy on test set: {:.3f}".format(svc.score(x_test_scaled, y_t
est)))
print()
print("Confusion Matrix ")
print(confusion_matrix(y_test, kn.predict(x_test)))
print()
print("Report")
print(classification_report(y_test, kn.predict(x_test)))

```

Accuracy on training set: 0.790

Accuracy on test set: 0.797

Confusion Matrix

```

[[105  20]
 [ 23  44]]

```

Report

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.82 | 0.84 | 0.83 | 125 |
| 1 | 0.69 | 0.66 | 0.67 | 67 |
| accuracy | | | 0.78 | 192 |
| macro avg | 0.75 | 0.75 | 0.75 | 192 |
| weighted avg | 0.77 | 0.78 | 0.77 | 192 |

C:\anaconda\lib\site-packages\sklearn\svm\base.py:193: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
 "avoid this warning.", FutureWarning)

PLOTTING BAR GRAPH BETWEEN DIFFERENT MODELS AND ACCURACY RESPECTIVELY

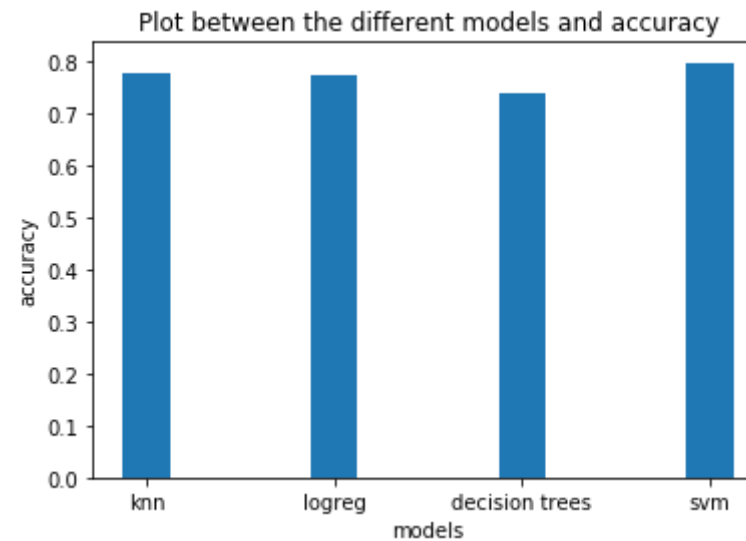
```

In [26]: list_of_models=['knn','logreg','decision trees','svm']
plt.bar(list_of_models,list_of_testing_accuracy,width = 0.25)
plt.xlabel('models')
plt.ylabel('accuracy')
plt.title('Plot between the different models and accuracy')

```

Out[26]: Text(0.5, 1.0, 'Plot between the different models and accuracy')

Out[26]: text(0.5, 1.0, 'Plot between the different models and accuracy')



From the above graph we can observe that the almost all the models have accuracy closely, comparing to Support Vector Machine has the high accuracy comparing to all the models

```
In [32]: table = pd.DataFrame({'Model_name':list_of_models, 'Accuracy':list_of_testing_accuracy})
print(table)
```

| | Model_name | Accuracy |
|---|----------------|----------|
| 0 | knn | 0.776042 |
| 1 | logreg | 0.770833 |
| 2 | decision trees | 0.739583 |
| 3 | svm | 0.796875 |

Conclusion: The diabetes data set is trained to different models, built the different models and calculated the accuracy of each model. Among all the models support vector machine has the highest accuracy.